

Complete Working Example: Celery + Redis + MongoDB Join Pipeline

This document presents a fully self-contained example demonstrating how a Celery client enqueues a request into Redis, how a Celery worker performs a MongoDB join operation, and how the final result is stored back into Redis as a key–value entry.

1. Architecture Overview

The system follows an asynchronous request–processing model. Redis acts as the message broker, Celery workers execute background jobs, MongoDB performs native joins, and Redis also serves as an application-level cache for results.

- Client submits request using Celery `delay()`
- Redis stores the task in a queue
- Celery worker retrieves and executes the task
- Result stored in Redis key–value store

2. MongoDB Collections

Users collection document:

```
{ "_id": "u1", "name": "Alice", "email": "alice@example.com" }
```

Orders collection document:

```
{ "_id": "o101", "userId": "u1", "amount": 1200, "status": "PAID" }
```

3. Celery Configuration (celery_app.py)

Celery is configured with Redis as both broker and result backend. JSON serialization and late acknowledgements are enabled.

```
from celery import Celery

celery_app = Celery(
    "join_tasks",
    broker="redis://localhost:6379/0",
    backend="redis://localhost:6379/1",
)

celery_app.conf.update(
    task_serializer="json",
    result_serializer="json",
    accept_content=["json"],
    task_acks_late=True,
    worker_prefetch_multiplier=1,
)
```

4. Celery Worker Task (tasks.py)

The worker executes a MongoDB aggregation pipeline with \$lookup and stores the result in Redis as a key-value entry.

```
from celery_app import celery_app
from pymongo import MongoClient
import redis, json

mongo = MongoClient("mongodb://localhost:27017")
db = mongo.shop

redis_client = redis.Redis(
    host="localhost",
    port=6379,
    db=2,
    decode_responses=True
)

@celery_app.task(bind=True)
def join_users_orders(self, user_id):

    pipeline = [
        {"$match": {"_id": user_id}},
        {
            "$lookup": {
                "from": "orders",
                "localField": "_id",
                "foreignField": "userId",
                "as": "orders",
            }
        },
    ]

    result = list(db.users.aggregate(pipeline))

    redis_key = f"user_order_result:{user_id}"
    redis_client.set(redis_key, json.dumps(result), ex=300)

    return {
        "redis_key": redis_key,
        "count": len(result),
    }
```

5. Celery Client (client.py)

The client submits the asynchronous request without blocking.

```
from tasks import join_users_orders

task = join_users_orders.delay("u1")
print("Submitted task ID:", task.id)
```

6. Worker Startup

The worker is started using the Celery command-line interface.

```
celery -A celery_app worker -l info
```

7. Redis Key–Value Result

The joined data is stored in Redis as JSON. This storage is independent of the Celery result backend.

Key: `user_order_result:u1`

Value:

```
[
  {
    "_id": "u1",
    "name": "Alice",
    "email": "alice@example.com",
    "orders": [
      {
        "_id": "o101",
        "amount": 1200,
        "status": "PAID"
      }
    ]
  }
]
```

8. Client Result Retrieval (fetch_result.py)

The client retrieves the data directly from Redis.

```
import redis, json

r = redis.Redis(
    host="localhost",
    port=6379,
    db=2,
    decode_responses=True
)

data = r.get("user_order_result:u1")
print(json.loads(data))
```

9. Design Notes

- Redis acts as both broker and cache
- MongoDB performs the join operation
- Tasks must be idempotent
- TTL prevents stale cache buildup